

Implement Gibbs Sampler

Problem Statement (1):

Gibbs Sampler is an algorithm that is used to find a motif in a DNA, RNA, or protein sequences. A motif is a short, recurring pattern (nucleotide sequence/amino acids) in a DNA sequence and uncovering them is important because they often have biological significance.

Input(s): Integers k (the length of motifs we are looking for), t (the number of DNA sequences in the data), and N (number of iterations our algorithm will perform), followed by a collection of strings Dna (which of course represents the DNA sequence we are looking through for motifs).

Output: “The strings $BestMotifs$ resulting from running $GibbsSampler(Dna, k, t, N)$ with 20 random starts.” $BestMotifs$ is a set of recurring patterns within the given DNA sequence, a List (Python) of strings

Function(s):

Constraints:

Algorithm description/Pseudocode (2.1):

```
GIBBSAMPLER( $Dna, k, t, N$ )
  randomly select  $k$ -mers  $Motifs = (Motif_1, \dots, Motif_t)$  in each string
  from  $Dna$ 
   $BestMotifs \leftarrow Motifs$ 
  for  $j \leftarrow 1$  to  $N$ 
     $i \leftarrow Random(t)$ 
     $Profile \leftarrow$  profile matrix constructed from all strings in  $Motifs$ 
    except for  $Motif_i$ 
     $Motif_i \leftarrow Profile$ -randomly generated  $k$ -mer in the  $i$ -th sequence
    if  $Score(Motifs) < Score(BestMotifs)$ 
       $BestMotifs \leftarrow Motifs$ 
  return  $BestMotifs$ 
```

NOTE: This is pseudo-code from Rosalind's site but the explanation is my own 😊

Explanation (2.2)

We start the algorithm by randomly selecting k -mers to comprise our initial motifs ($BestMotifs$)

For N iterations we randomly select an index between 1 & t

We start constructing our profile matrix based on all the motifs except the one at index i

This matrix represents the probabilities of each nucleotide at each position in the motif

We then use the profile matrix to randomly generate a k -mer in the i -th DNA sequence.

With this profile matrix, we can calculate a score to evaluate how well the current set of motifs compares to the DNA sequence provided.

If the score of the current motifs is higher than $bestMotifs$, we update $bestMotifs$ to the current one

After N iterations, (Rosalind asks for 20 runs), we should have the 'best' set of motifs found.

Time analysis (3):

This problem has a run-time complexity of $O(t * k^2 + N * (n - k + 1) * k^2)$ for a single iteration of the algorithm

This problem has a run-time complexity of $O(N * t * k^2 + N * (n - k + 1) * k^2)$ for a single iteration of the algorithm

N is the number of iterations for which we run Gibbs Sampler

n is the length of the given DNA sequence

Step by Step:

Randomly selecting k-mers requires us to run windows of k length through t DNA sequences $\Rightarrow O(t * k)$

Creating the profile matrix requires us to iterate over the motifs t times, and for each motif we iterate k

Iterations AND calculating probabilities at each position which takes $O(k)$ as well $\Rightarrow O(t * k * k) \Rightarrow O(t * k^2)$

Updating motif _{i} requires us to replace one k-mer with another, taking another $O(k)$

And finally comparing the scores and updating the current bestMotifs is done in constant time $O(1)$

After incorporating the number of times we iterate for a k-mer $(n - k + 1)$ and factoring the expression: we arrive at

$O(t * k^2 + N * (n - k + 1) * k^2)$ for a single iteration of the algorithm.

Discussion (5):

The approach for this algorithm was a bit unique.

After looking for some guidance for how to complete this algorithm online, I found that this could be better visualized in a Jupyter notebook, which would have been my preference as there are a lot of cool statistical libraries we could have used to automate and perhaps even optimize the runtime of our algorithm.

The runtime itself is a bit concerning as this algorithm will take a considerable amount of time on bigger k-mers and DNA sequences.

As usual when it comes to pattern-finding algorithms, we can always try and utilize machine learning in interesting ways I have not explored on my own yet. While this may be more computationally expensive, the potential accuracies and discoveries could be limitless.

Some edge cases may cause a high N number of iterations before coming upon a viable set of motifs.